

# DAM lecture 8:

Dimensionality reduction 2

03.03.2016

Aasa Feragen

aasa@diku.dk

# Assignment 1 feedback

- ▶ Feedback is now available
- ▶ We have been strict on presentation – most of you will have an easy update
- ▶ Use the TA sessions!
- ▶ Questions?

## After today's lecture you should

- ▶ be familiar with the equivalent definitions of PCA by least squares projection error minimization, projected variance maximization, low distortion embedding, and eigenvalue decomposition of the covariance matrix
- ▶ be able to interpret the different equivalent PCA definitions and use them to pinpoint strengths and weaknesses of PCA
- ▶ be able to use PCA for visualization of global dataset variation
- ▶ be familiar with the curse of dimensionality and the need for dimensionality reduction
- ▶ be familiar with partial derivatives, gradients, and their use for finding principal components (if time allows)

# Literature for today's lecture

- ▶ Chapters 4 and 10
- ▶ **Shlens tutorial:**  
Optional; fantastic intro to PCA with Matlab code  
(find it on Absalon)


## Recall from Lecture 6: Covariance matrix

- ▶ For a sampled dataset  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\} \subset \mathbb{R}^d$ , we can define its  $d \times d$  *covariance matrix*  $\Sigma$  by setting

$$\Sigma_{i,j} = \text{cov}(x_i, x_j).$$

- ▶ The variance of each coordinate is found along the diagonal:  
 $s_{x_i}^2 = \lambda_i$

## Recall from Lecture 6: Decomposing the covariance matrix



### Theorem (Eigenvalue decomposition)

If  $\Sigma$  is a  $d \times d$  matrix with linearly independent eigenvectors  $\mathbf{e}_1, \dots, \mathbf{e}_d$ , with corresponding eigenvalues  $\lambda_1, \dots, \lambda_d$ , then  $\Sigma$  has a decomposition

$$\Sigma = Q \begin{pmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \lambda_d \end{pmatrix} Q^{-1},$$

where the columns of  $Q$  are the eigenvectors  $\mathbf{e}_1, \dots, \mathbf{e}_d$ .


## Recall from Lecture 6: What does the eigenvalue decomposition of the covariance matrix mean?

- What does

$$\Sigma = Q \begin{pmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \lambda_d \end{pmatrix} Q^{-1}$$

mean?

- $Q$  is a change of bases, re-expressing the covariance matrix in the basis defined by the eigenvectors.




## Recall from Lecture 6: What does the eigenvalue decomposition of the covariance matrix mean?

- ▶ The diagonal matrix

$$D = \begin{pmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \lambda_d \end{pmatrix}$$

is the covariance of the dataset  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\} \subset \mathbb{R}^d$ , expressed in the new basis.


- ▶ What do you see?
  - ▶ The coordinates of the data points in the new basis are independent!
  - ▶ The variance of each coordinate is found along the diagonal!





## Recall from Lecture 6: What does the change of basis do?


- ▶ Align principal components with axes in the new coordinate system
- ▶ The intrinsic geometry of the data is unchanged! Only rotation and reflection.  
(Because eigenvectors are orthonormal)



## Recall from Lecture 6: Principal components analysis (PCA)

The  $k$  first *principal components* of the dataset  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$  span the  $k$ -dimensional linear subspace  $V \subset \mathbb{R}^d$  that *maximizes the variance* of the projected dataset

- **Question:** Is this  $V$  unique?




# Dimensionality reduction

- ▶ PCA is an example of *dimensionality reduction*
- ▶ Dimensionality reduction refers to the process of reducing the dimensionality in your data representation.
- ▶ More precisely: Given a dataset  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\} \subset \mathbb{R}^{d_1}$ , finding a representation of your dataset

$$\{\phi(\mathbf{x}_1), \phi(\mathbf{x}_2), \dots, \phi(\mathbf{x}_N)\} \subset \mathbb{R}^{d_2}$$


where  $d_2 < d_1$ , and where you retain the properties of your dataset as well as possible.

- ▶ Why is this useful?



# The curse of dimensionality<sup>1</sup>

- ▶ In order to sample the interval  $[0, 1]$  with density 0.1, I need 10 points.
- ▶ In order to sample the cube  $[0, 1] \times [0, 1]$  with the same density, I need 100 points.
- ▶ etc
- ▶ The more dimensions, the more data you need for drawing conclusions.



---




<sup>1</sup>Figure from Bishop: Pattern Recognition and Machine Learning

# The curse of dimensionality<sup>1</sup>

- ▶ Consider the  $d$ -cube  $[-1, 1]^d$ .
- ▶ The distance from the center to a corner is

$$\sqrt{d} \rightarrow \infty \text{ as } d \rightarrow \infty$$

- ▶ When  $d$  gets large, everything gets large – including noise effects!






<sup>1</sup>Figure from Bishop: Pattern Recognition and Machine Learning

# The curse of dimensionality<sup>1</sup>

- ▶ Consider the  $d$ -cube  $[-1, 1]^d$ .
- ▶ The distance from the center to a corner is

$$\sqrt{d} \rightarrow \infty \text{ as } d \rightarrow \infty$$

- ▶ When  $d$  gets large, everything gets large – including noise effects!
- ▶ What sort of problems could this give you?




---

<sup>1</sup>Figure from Bishop: Pattern Recognition and Machine Learning


# Equivalence of projection error minimization / projected variance maximization

- ▶ PCA equivalently formulated as minimizing squared projection error
- ▶ A least squares problem
- ▶ Equivalent to variance maximization *up to projection*



# Equivalence of projection error minimization / projected variance maximization


- ▶ PCA equivalently formulated as minimizing squared projection error
- ▶ A least squares problem
- ▶ Equivalent to variance maximization *up to projection* (**is it?**)






# Equivalence of projection error minimization / projected variance maximization

- ▶ PCA equivalently formulated as minimizing squared projection error
- ▶ A least squares problem
- ▶ Equivalent to variance maximization *up to projection* (**is it?**)
- ▶ For equivalence: Ask PCs to pass through the mean





# Equivalence of projection error minimization / projected variance maximization

- ▶ PCA equivalently formulated as minimizing squared projection error
- ▶ A least squares problem
- ▶ Equivalent to variance maximization *up to projection* (**is it?**) (**why equivalent?**)
- ▶ For equivalence: Ask PCs to pass through the mean





# PCA and low distortion embedding – Multidimensional Scaling

- Projection of data onto PCs is often used to visualize global dataset structure – below is the talking face dataset



# PCA and low distortion embedding – Multidimensional Scaling

- ▶ Projection of data onto PCs is often used to visualize global dataset structure – below is the talking face dataset






- ▶ PCA defines minimizing subspaces for the objective function

$$\sum_{i,j=1}^N (\|\mathbf{x}_i - \mathbf{x}_j\|_{\mathbb{R}^d}^2 - \|\text{pr}_V(\mathbf{x}_i) - \text{pr}_V(\mathbf{x}_j)\|_V^2)$$

- ▶ Does PCA define a unique minimizing subspace?

# PCA and low distortion embedding – Multidimensional Scaling

- ▶ Projection of data onto PCs is often used to visualize global dataset structure – below is the talking face dataset



- ▶ PCA defines minimizing subspaces for the objective function

$$\sum_{i,j=1}^N (\|\mathbf{x}_i - \mathbf{x}_j\|_{\mathbb{R}^d}^2 - \|\text{pr}_V(\mathbf{x}_i) - \text{pr}_V(\mathbf{x}_j)\|_V^2)$$

- ▶ Does PCA define a unique minimizing subspace?
- ▶ The strategy of projecting onto subspaces while preserving distances is called *multidimensional scaling*, or MDS

# Projection onto a subspace


- **Task:** Project your datapoints  $\mathbf{x}_i$  onto the linear subspace  $V$  spanned by  $\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_k$ .

# Projection onto a subspace

- ▶ **Task:** Project your datapoints  $\mathbf{x}_i$  onto the linear subspace  $V$  spanned by  $\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_k$ .
- ▶ **How do you do that? (You will be doing this in a future assignment)**

# Projection onto a subspace

- ▶ **Task:** Project your datapoints  $\mathbf{x}_i$  onto the linear subspace  $V$  spanned by  $\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_k$ .
- ▶ **How do you do that? (You will be doing this in a future assignment)**
- ▶  $\text{pr}_{\mathbf{v}}(\mathbf{w}) = \mathbf{w} \cdot \mathbf{v} = \mathbf{w}^T \mathbf{v}$  for a unit vector  $\mathbf{v}$  – given by dot product




$$\text{pr}_{\mathbf{v}} \mathbf{w} = \mathbf{v} \cdot \mathbf{w} = \mathbf{v}^T \mathbf{w} = \mathbf{w}^T \mathbf{v} \text{ if } \|\mathbf{v}\| = 1$$



# Projection onto a subspace

- ▶ **Task:** Project your datapoints  $\mathbf{x}_i$  onto the linear subspace  $V$  spanned by  $\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_k$ .
- ▶ **How do you do that? (You will be doing this in a future assignment)**
- ▶  $\text{pr}_V(\mathbf{w}) = \mathbf{w} \cdot \mathbf{v} = \mathbf{w}^T \mathbf{v}$  for a unit vector  $\mathbf{v}$  – given by dot product
- ▶  $\text{pr}_V(\mathbf{w}) = (\mathbf{w}^T \mathbf{e}_1)\mathbf{e}_1 + \dots + (\mathbf{w}^T \mathbf{e}_k)\mathbf{e}_k = (\mathbf{w}^T \mathbf{e}_1, \dots, \mathbf{w}^T \mathbf{e}_k)$




$$\text{pr}_V \mathbf{w} = \mathbf{v} \cdot \mathbf{w} = \mathbf{v}^T \mathbf{w} = \mathbf{w}^T \mathbf{v} \text{ if } \|\mathbf{v}\| = 1$$

# Why does PCA give MDS?

**Input:** Distance matrix  $D = (d_{ij})$

for distances  $d_{ij} = d(z_i, z_j)$

dataset  $\{z_n\}_{n=1}^N \subset X$  general data space



# Why does PCA give MDS?

**Input:** Distance matrix  $D = (d_{ij})$


for distances  $d_{ij} = d(z_i, z_j)$

dataset  $\{z_n\}_{n=1}^N \subset X$  general data space

**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.



# Why does PCA give MDS?

**Input:** Distance matrix  $D = (d_{ij})$

for distances  $d_{ij} = d(z_i, z_j)$


dataset  $\{z_n\}_{n=1}^N \subset X$  general data space

**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

That is, distances are preserved as well as possible.



# Why does PCA give MDS?

**Input:** Distance matrix  $D = (d_{ij})$

for distances  $d_{ij} = d(z_i, z_j)$

dataset  $\{z_n\}_{n=1}^N \subset X$  general data space


**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

That is, distances are preserved as well as possible.

**Assume**  $X = \mathbb{R}^k$  for  $k \gg 0$



# Why does PCA give MDS?

**Input:** Distance matrix  $D = (d_{ij})$

for distances  $d_{ij} = d(z_i, z_j)$

dataset  $\{z_n\}_{n=1}^N \subset X$  general data space

**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that


$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

That is, distances are preserved as well as possible.

**Assume**  $X = \mathbb{R}^k$  for  $k \gg 0$

$f(z) = Mz$  linear projection onto linear subspace  $\subset \mathbb{R}^k$   
of low dimension.



# Why does PCA give MDS?


**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

That is, we seek

$$\operatorname{argmin} \sum_{i,j} (d_{ij}^2 - \|Mz_i - Mz_j\|^2)$$



# Why does PCA give MDS?


**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

That is, we seek

$$\begin{aligned} & \operatorname{argmin} \sum_{i,j} (d_{ij}^2 - \|Mz_i - Mz_j\|^2) \\ &= \operatorname{argmin} \left( \sum_{i,j} d_{ij}^2 - \sum_{i,j} \|Mz_i - Mz_j\|^2 \right) \end{aligned}$$





# Why does PCA give MDS?


**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

That is, we seek

$$\begin{aligned} & \operatorname{argmin} \sum_{i,j} (d_{ij}^2 - \|Mz_i - Mz_j\|^2) \\ &= \operatorname{argmin} \left( \sum_{i,j} d_{ij}^2 - \sum_{i,j} \|Mz_i - Mz_j\|^2 \right) \\ &= \operatorname{argmax} \sum_{i,j} \|Mz_i - Mz_j\|^2 \end{aligned}$$



# Why does PCA give MDS?


**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

That is, we seek

$$\begin{aligned} & \operatorname{argmin} \sum_{i,j} (d_{ij}^2 - \|Mz_i - Mz_j\|^2) \\ &= \operatorname{argmin} \left( \sum_{i,j} d_{ij}^2 - \sum_{i,j} \|Mz_i - Mz_j\|^2 \right) \\ &= \operatorname{argmax} \sum_{i,j} \|Mz_i - Mz_j\|^2 \\ &= \operatorname{argmax} 2N \sum_i \|Mz_i - \frac{1}{N} \sum_j Mz_j\|^2 \end{aligned}$$



# Why does PCA give MDS?


**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

That is, we seek

$$\begin{aligned} & \operatorname{argmin} \sum_{i,j} (d_{ij}^2 - \|Mz_i - Mz_j\|^2) \\ &= \operatorname{argmin} \left( \sum_{i,j} d_{ij}^2 - \sum_{i,j} \|Mz_i - Mz_j\|^2 \right) \\ &= \operatorname{argmax} \sum_{i,j} \|Mz_i - Mz_j\|^2 \\ &= \operatorname{argmax} 2N \sum_i \|Mz_i - \frac{1}{N} \sum_j Mz_j\|^2 \\ & \text{equivalent to maximizing projected variance.} \end{aligned}$$



# Why does PCA give MDS?

**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$


is minimized.

That is, we seek

$$\begin{aligned} & \operatorname{argmin} \sum_{i,j} (d_{ij}^2 - \|Mz_i - Mz_j\|^2) \\ &= \operatorname{argmin} \left( \sum_{i,j} d_{ij}^2 - \sum_{i,j} \|Mz_i - Mz_j\|^2 \right) \\ &= \operatorname{argmax} \sum_{i,j} \|Mz_i - Mz_j\|^2 \\ &= \operatorname{argmax} 2N \sum_i \|Mz_i - \frac{1}{N} \sum_j Mz_j\|^2 \end{aligned}$$

equivalent to maximizing projected variance.

Familiar?



# Why does PCA give MDS?

**Input:** Distance matrix  $D = (d_{ij})$

for distances  $d_{ij} = d(z_i, z_j)$


dataset  $\{z_n\}_{n=1}^N \subset X$  general data space

**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that

$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

**New interpretation of PCA:** Preserving squared distances



# Why does PCA give MDS?

**Input:** Distance matrix  $D = (d_{ij})$

for distances  $d_{ij} = d(z_i, z_j)$

dataset  $\{z_n\}_{n=1}^N \subset X$  general data space

**Goal:** Find mapping  $f: X \rightarrow \mathbb{R}^d$  for small  $d$  such that


$$\Phi(Y) = \sum_{i=1}^N \sum_{j=1}^N (d_{ij}^2 - \|f(z_i) - f(z_j)\|^2)$$

is minimized.

**New interpretation of PCA:** Preserving squared distances





**Revealing problem with PCA:**

Preserves long distances better than short ones (why problem?)



# PCA and visualization

- Low-dimensional representation of the data via projection onto first 2 or 3 principal components



# Summary

We have seen two alternative ways of defining PCA:

- ▶ Minimizing projection error

$$\operatorname{argmin}_V \sum_{n=1}^N \|\mathbf{x}_n - \operatorname{pr}_V(\mathbf{x}_n)\|^2$$

- ▶ Minimizing distortion

$$\operatorname{argmin}_V \sum_{i,j=1}^N (\|\mathbf{x}_i - \mathbf{x}_j\|_{\mathbb{R}^d}^2 - \|\operatorname{pr}_V(\mathbf{x}_i) - \operatorname{pr}_V(\mathbf{x}_j)\|_V^2)$$

- ▶ What do we learn from these?



# Summary

We have seen two alternative ways of defining PCA:

- ▶ Minimizing projection error

$$\operatorname{argmin}_V \sum_{n=1}^N \|\mathbf{x}_n - \operatorname{pr}_V(\mathbf{x}_n)\|^2$$


- ▶ Minimizing distortion

$$\operatorname{argmin}_V \sum_{i,j=1}^N (\|\mathbf{x}_i - \mathbf{x}_j\|_{\mathbb{R}^d}^2 - \|\operatorname{pr}_V(\mathbf{x}_i) - \operatorname{pr}_V(\mathbf{x}_j)\|_V^2)$$

- ▶ What do we learn from these? PCA is sensitive to outliers!

## Alternate way of solving PCA: Optimization

The PC1 of the dataset  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$  spans the line  $l \subset \mathbb{R}^d$  that *maximizes the variance* of the projected dataset




Assuming  $\bar{\mathbf{x}} = 0$ , this translates to the optimization problem

$$\begin{aligned} & \operatorname{argmin}_{\mathbf{w}, \|\mathbf{w}\|=1} \sum_{n=1}^N \|\operatorname{pr}_{\mathbf{w}} \mathbf{x}_n - \overline{\operatorname{pr}_{\mathbf{w}}(\mathbf{x})}\|^2 \\ &= \operatorname{argmin}_{\mathbf{w}, \|\mathbf{w}\|=1} \sum_{n=1}^N \|\mathbf{w}^T \mathbf{x}_n\|^2, \end{aligned}$$

where  $\mathbf{w}$  is a unit vector parallel to the line  $l$ .

# Optimization: Gradient descent

- ▶ We will be using gradient descent intensively in the next couple of weeks, so let's start out reviewing its basics
- ▶ Gradient descent is used to find *minimum values*
- ▶ Idea: Walk downhill until you hit bottom




Issues that need your attention:

- ▶ Where do I start?
- ▶ How long steps?
- ▶ When have I hit bottom?
- ▶ Do you see any dangers?

# Optimization: Gradient descent

- ▶ We will be using gradient descent intensively in the next couple of weeks, so let's start out reviewing its basics
- ▶ Gradient descent is used to find *minimum values*
- ▶ Idea: Walk downhill until you hit bottom




Issues that need your attention:

- ▶ Where do I start?  
Common: Initialize Random sample from normal distribution
- ▶ How long steps?
- ▶ When have I hit bottom?
- ▶ Do you see any dangers?

# Optimization: Gradient descent

- ▶ We will be using gradient descent intensively in the next couple of weeks, so let's start out reviewing its basics
- ▶ Gradient descent is used to find *minimum values*
- ▶ Idea: Walk downhill until you hit bottom




Issues that need your attention:

- ▶ Where do I start?  
Common: Initialize Random sample from normal distribution
- ▶ How long steps? Common:  $\eta_t = \eta \|\nabla E\|$ ; typical  $\nabla = 0.1$
- ▶ When have I hit bottom?
- ▶ Do you see any dangers?

# Optimization: Gradient descent

- ▶ We will be using gradient descent intensively in the next couple of weeks, so let's start out reviewing its basics
- ▶ Gradient descent is used to find *minimum values*
- ▶ Idea: Walk downhill until you hit bottom



Issues that need your attention:

- ▶ Where do I start?  
Common: Initialize Random sample from normal distribution
- ▶ How long steps? Common:  $\eta_t = \eta \|\nabla E\|$ ; typical  $\nabla = 0.1$
- ▶ When have I hit bottom? Threshold on # steps or step size
- ▶ Do you see any dangers?

## Before gradient descent: Derivatives

- ▶ Let  $f: \mathbb{R} \rightarrow \mathbb{R}$ , given by  $y = f(x)$ , be a function
- ▶ The *derivative* of  $f$  in the point  $x_0$  is the limit

$$\frac{df}{dx}(x_0) = f'(x_0) = \lim_{h \rightarrow 0} \frac{f(x_0 + h) - f(x_0)}{h}.$$

- ▶ When the derivative exists in every  $x_0$ , the function  $f$  is *differentiable*.
- ▶ Using the theory of limits, one can derive formulas for common derivatives:

$$\begin{aligned} \frac{d}{dx} x^n &= n x^{n-1}, & \frac{d}{dx} (f(g(x))) &= f'(g(x)) g'(x) \\ \frac{d}{dx} (f(x) + g(x)) &= \frac{d}{dx} f(x) + \frac{d}{dx} g(x) \end{aligned}$$

## Before gradient descent: Derivatives

- ▶ Let  $f: \mathbb{R} \rightarrow \mathbb{R}$ , given by  $y = f(x)$ , be a function
- ▶ The *derivative* of  $f$  in the point  $x_0$  is the limit

$$\frac{df}{dx}(x_0) = f'(x_0) = \lim_{h \rightarrow 0} \frac{f(x_0 + h) - f(x_0)}{h}.$$

- ▶ When the derivative exists in every  $x_0$ , the function  $f$  is *differentiable*.
- ▶ Using the theory of limits, one can derive formulas for common derivatives:

$$\begin{aligned} \frac{d}{dx} x^n &= n x^{n-1}, & \frac{d}{dx} (f(g(x))) &= f'(g(x)) g'(x) \\ \frac{d}{dx} (f(x) + g(x)) &= \frac{d}{dx} f(x) + \frac{d}{dx} g(x) \end{aligned}$$

- ▶ What are the following derivatives?


$$\frac{d}{dx} (a - x)^2 \qquad \frac{d}{dx} \sum_{n=1}^N (x_n - x)^2$$



# Partial derivatives – informally

## Example


The function  $f(x, y) = x^2 y^3$ .



# Partial derivatives – informally

## Example

The function  $f(x, y) = x^2 y^3$ .




$$\frac{\partial}{\partial x} f(x, y) = 2xy^3 \quad \frac{\partial}{\partial y} f(x, y) = x^2 3y^2$$

# Partial derivatives – informally

## Example

The function  $f(x, y) = x^2 y^3$ .



$$\frac{\partial}{\partial x} f(x, y) = 2xy^3 \quad \frac{\partial}{\partial y} f(x, y) = x^2 3y^2$$

## Exercise

What are  $\frac{\partial}{\partial x}$  and  $\frac{\partial}{\partial y}$  for the following functions?

- ▶  $f(x, y) = y^2 \sin^2(x)$
- ▶  $f(x, y) = \sin(x)e^{yx}$

# Partial derivatives

Recall that the derivative of a function  $f: \mathbb{R} \rightarrow \mathbb{R}$  at  $x_0 \in \mathbb{R}$  is given by

$$\frac{\partial f}{\partial x}(x_0) = \lim_{h \rightarrow 0} \frac{f(x_0 + h) - f(x_0)}{h}$$

Let  $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ ,  $z = f(x, y)$  be a function of two variables.

- ▶ The partial derivative of  $f$  with respect to  $x$  at  $(x_0, y_0)$  is

$$\frac{\partial f}{\partial x}(x_0, y_0) = \lim_{h \rightarrow 0} \frac{f(x_0 + h, y_0) - f(x_0, y_0)}{h}$$

# Partial derivatives

Recall that the derivative of a function  $f: \mathbb{R} \rightarrow \mathbb{R}$  at  $x_0 \in \mathbb{R}$  is given by

$$\frac{\partial f}{\partial x}(x_0) = \lim_{h \rightarrow 0} \frac{f(x_0 + h) - f(x_0)}{h}$$

Let  $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ ,  $z = f(x, y)$  be a function of two variables.

- ▶ The partial derivative of  $f$  with respect to  $x$  at  $(x_0, y_0)$  is

$$\frac{\partial f}{\partial x}(x_0, y_0) = \lim_{h \rightarrow 0} \frac{f(x_0 + h, y_0) - f(x_0, y_0)}{h}$$

- ▶ The partial derivative of  $f$  with respect to  $y$  at  $(x_0, y_0)$  is

$$\frac{\partial f}{\partial y}(x_0, y_0) = \lim_{h \rightarrow 0} \frac{f(x_0, y_0 + h) - f(x_0, y_0)}{h}$$

# Partial derivatives

- ▶  $f$  is *differentiable* at  $(x_0, y_0)$  if both partial derivatives exist there

# Partial derivatives

- ▶  $f$  is *differentiable* at  $(x_0, y_0)$  if both partial derivatives exist there
- ▶  $f$  is *differentiable* if it is differentiable everywhere

# Partial derivatives

- ▶  $f$  is *differentiable* at  $(x_0, y_0)$  if both partial derivatives exist there
- ▶  $f$  is *differentiable* if it is differentiable everywhere
- ▶ Concept carries over to functions  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  of any number of variables.



# Gradients

The *gradient* of a function  $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ ,  $z = f(x, y)$  in the point  $(x_0, y_0)$  is given by


$$\begin{aligned}\nabla f(x_0, y_0) &= \left( \frac{\partial f}{\partial x}(x_0, y_0), \frac{\partial f}{\partial y}(x_0, y_0) \right) \\ &= \lim_{h \rightarrow 0} \left( \frac{f(x_0+h, y_0) - f(x_0, y_0)}{h}, \frac{f(x_0, y_0+h) - f(x_0, y_0)}{h} \right)\end{aligned}$$

# Gradients

The *gradient* of a function  $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ ,  $z = f(x, y)$  in the point  $(x_0, y_0)$  is given by

$$\begin{aligned}\nabla f(x_0, y_0) &= \left( \frac{\partial f}{\partial x}(x_0, y_0), \frac{\partial f}{\partial y}(x_0, y_0) \right) \\ &= \lim_{h \rightarrow 0} \left( \frac{f(x_0+h, y_0) - f(x_0, y_0)}{h}, \frac{f(x_0, y_0+h) - f(x_0, y_0)}{h} \right)\end{aligned}$$

The gradient points in the direction of the steepest descent of the function.




# Gradients

The *gradient* of a function  $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ ,  $z = f(x, y)$  in the point  $(x_0, y_0)$  is given by

$$\begin{aligned}\nabla f(x_0, y_0) &= \left( \frac{\partial f}{\partial x}(x_0, y_0), \frac{\partial f}{\partial y}(x_0, y_0) \right) \\ &= \lim_{h \rightarrow 0} \left( \frac{f(x_0+h, y_0) - f(x_0, y_0)}{h}, \frac{f(x_0, y_0+h) - f(x_0, y_0)}{h} \right)\end{aligned}$$

The gradient points in the direction of the steepest descent of the function.




What can this be useful for?

# Gradients

The *gradient* of a function  $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ ,  $z = f(x, y)$  in the point  $(x_0, y_0)$  is given by

$$\begin{aligned}\nabla f(x_0, y_0) &= \left( \frac{\partial f}{\partial x}(x_0, y_0), \frac{\partial f}{\partial y}(x_0, y_0) \right) \\ &= \lim_{h \rightarrow 0} \left( \frac{f(x_0+h, y_0) - f(x_0, y_0)}{h}, \frac{f(x_0, y_0+h) - f(x_0, y_0)}{h} \right)\end{aligned}$$


The gradient points in the direction of the steepest descent of the function.



What can this be useful for? **Optimization – gradient descent!**

# Optimization: Gradient descent

- ▶ We will be using gradient descent intensively in the next couple of weeks, so let's start out reviewing its basics
- ▶ Gradient descent is used to find *minimum values*
- ▶ Idea: Walk downhill until you hit bottom



Issues that need your attention:

- ▶ Where do I start?
- ▶ How long steps?
- ▶ When have I hit bottom?
- ▶ Do you see any dangers?

# Basic gradient descent algorithm for PCA

Searching for

$$\operatorname{argmin}_{\mathbf{w}, \|\mathbf{w}\|=1} f(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}, \|\mathbf{w}\|=1} \sum_{n=1}^N \|\mathbf{w}^T \mathbf{x}_n\|^2$$

- 1: Initialize:  $\mathbf{w} \leftarrow \mathbf{w}_0$
- 2: **while** Convergence = False **do**
- 3:   Compute  $\nabla f(\mathbf{w}_0)$
- 4:   Set step size  $s$
- 5:    $\mathbf{w} \leftarrow \mathbf{w} + s \nabla f(\mathbf{w}_0)$
- 6:    $\mathbf{w} \leftarrow \frac{\mathbf{w}}{\|\mathbf{w}\|}$
- 7:   **if** Convergence criteria hold **then**
- 8:     Convergence = True
- 9:   **end if**
- 10: **end while**

# Basic gradient descent algorithm for PCA

Searching for

$$\operatorname{argmin}_{\mathbf{w}, \|\mathbf{w}\|=1} f(\mathbf{w}) = \operatorname{argmin}_{\mathbf{w}, \|\mathbf{w}\|=1} \sum_{n=1}^N \|\mathbf{w}^T \mathbf{x}_n\|^2$$

- 1: Initialize:  $\mathbf{w} \leftarrow \mathbf{w}_0$
- 2: **while** Convergence = False **do**
- 3:   Compute  $\nabla f(\mathbf{w}_0)$
- 4:   Set step size  $s$
- 5:    $\mathbf{w} \leftarrow \mathbf{w} + s \nabla f(\mathbf{w}_0)$
- 6:    $\mathbf{w} \leftarrow \frac{\mathbf{w}}{\|\mathbf{w}\|}$  *Why do we need this?*
- 7:   **if** Convergence criteria hold **then**
- 8:     Convergence = True
- 9:   **end if**
- 10: **end while**

# When is gradient descent useful?

- ▶ For PCA? Large, high-dimensional datasets
- ▶ For other things? We shall see next week!



## Next week:

- ▶ Tuesday: Regression with a single variable, and a bit of linear algebra (Chapter 4 + 14)
- ▶ Thursday: Multivariate regression (Chapter 15)